CAI Assignment - II

PS: Comparative Financial QA System: RAG vs Fine-Tuning

Group Number: 71

Name	BITS ID	Contribution
Kulkarni Siddharth Prasad	2023AC05082	100%
Patil Omkar Mahesh	2023AC05085	100%
Dhiman Kundu	2023AC05129	100%
Sarit Ghosh	2023AC05131	100%
Soumen Choudhury	2023AC05143	100%

Hosted APP Details:

- Due to the 1 GB storage limit on free platforms, we fine-tuned an additional smaller GPT-2 model (~400 MB). Its accuracy is low but was deployable. All reported metrics, are based on the fine-tuned GPT-2 Medium model (~1.4 GB).
- **GitHub Repo URL:** https://github.com/mykeysid10/Bits-CAI-G71-A2
- **GPT-2-Medium Finetuned Model URL:** Download the models.safetensors file in GUI/finetuned-gpt2m-artifacts/ folder: https://drive.google.com/file/d/17JbY8qm715LImyr1J7de-oxzGSHf-3JN/view?usp=sharing
- **APP URL:** https://bits-cai-group-71-assignment-2.streamlit.app/
- **RAW Dataset URL:** https://drive.google.com/drive/folders/1Lc62ezVYTq62yMrwCAmtkAJYf__qSFYK?usp=sharing

Contents

Part	Content	Page Number
I	Data Collection & Pre-processing.	2
II	Retrieval-Augmented Generation System Implementation.	3
III	Fine-Tuned Model System Implementation.	4
IV	Testing, Evaluation & Comparison	5
V	Code Structure	8
VI	UI Snapshots	9

Data Overview: Financial statements for last 2 years.

- **Company:** PHILLIPS EDISON & COMPANY, INC.
- **Source URLs:**
    - https://www.sec.gov/ix?doc=/Archives/edgar/data/0001476204/000147620425000065/peco-20250331.htm
    - https://www.sec.gov/Archives/edgar/data/1476204/000147620425000092/peco-20250630.htm

I. Data Collection & Pre-Processing

Steps Performed

- ✓ **Text Extraction:** Converted PDF financial statements to plain text using ***OCR***.
- ✓ **Data Cleaning & Summarization:** Applied Groq’s ***deepseek-r1-distill-llama-70b*** model to condense lengthy narratives, calculations, and insights into concise sentences.
- ✓ **Noise Removal & Segmentation:**
  - Removed headers, footers, and page numbers.
  - Structured cleaned data into well-defined sections (as shown below).
- ✓ **Q&A Generation:**
  - Created ***75 QnA pairs*** representing key facts from the financial data and stored it to ***financial_qna_pairs.csv***.

Segmented .txt File Structure

- Section 1: Income Statement
- Section 2: Balance Sheet
- Section 3: Cash Flow Statement
- Section 4: Real Estate Portfolio
- Section 5: Leasing & Occupancy
- Section 6: Debt & Financing
- Section 7: Equity & Distributions
- Section 8: Risk Factors & Strategic Outlook

Sample Q&A Output

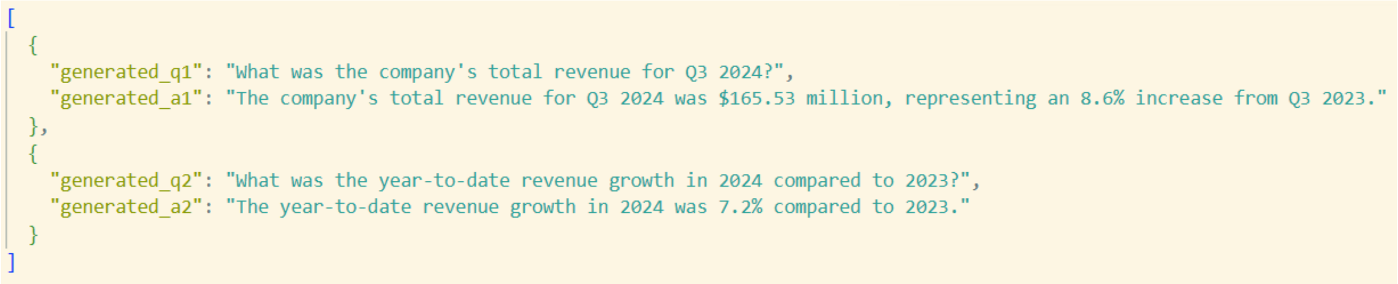

```
[
  {
    "generated_q1": "What was the company's total revenue for Q3 2024?",
    "generated_a1": "The company's total revenue for Q3 2024 was $165.53 million, representing an 8.6% increase from Q3 2023."
  },
  {
    "generated_q2": "What was the year-to-date revenue growth in 2024 compared to 2023?",
    "generated_a2": "The year-to-date revenue growth in 2024 was 7.2% compared to 2023."
  }
]
```

II. Retrieval-Augmented Generation System Implementation.

Steps Performed

- ✓ Used the same ***75*** question–answer pairs from the RAG pipeline and converted them into an instruction–response format for fine-tuning.
- ✓ Processed financial Q&A data by splitting answers into optimized chunks (100 & 400 tokens) with metadata (IDs, sources) for efficient retrieval.
- ✓ Built dual retrieval systems:
  - Dense retrieval using FAISS with *all-MiniLM-L6-v2* embeddings.
  - Sparse retrieval via BM25 with stop word/punctuation filtering.
- ✓ Implemented hybrid search to merge results from both retrievers, normalizing scores (1/(1+distance) for FAISS, raw BM25 scores) for balanced ranking.
- ✓ Added two-stage refinement:
  - Broad-phase: Hybrid retrieval fetches top-5 chunks.
  - Precision-phase: Re-ranks candidates using a cross-encoder (*ms-marco-MiniLM-L-6-v2*) for higher accuracy.
- ✓ **Generated answers** with *roberta-base-squad2*, truncating context to 1024 tokens to avoid overload.
- ✓ **Deployed guardrails** to block unsafe queries:
  - Violence (kill, attack → "Blocked: violence-related terms detected").
  - Financial crimes (fraud, scam → "Blocked: illegal activity query").
  - Personal data (SSN, password → "Blocked: sensitive info request").
  - Short queries (<5 chars → "Query too short").
- ✓ Integrated the ***RAG*** model into the existing user interface, ensuring guardrail checks run before model processing.
- ✓ Evaluated performance via:
  - Semantic similarity (cosine scores between generated/true answers).
  - Exact match rate for factual precision.

Conclusion

1. Successfully developed a complete pipeline integrating hybrid retrieval, reranking and QA generation, optimized for financial domain accuracy.
2. Achieved ***0.981*** average similarity score, demonstrating strong contextual comprehension beyond exact keyword matching.
3. Delivered consistent ***sub-0.5s*** response times with modular components for easy maintenance and upgrades.
4. Excelled in numeric precision ***(e.g., "$161.78M")***, query safety, and transparent sourcing—key for financial applications.
5. Focus areas include confidence score calibration, growth rate notation handling and edge-case testing for enterprise readiness.

III. Fine-Tuned Model System Implementation.

Steps Performed

- ✓ Used the same ***75*** question–answer pairs from the RAG pipeline and converted them into an instruction–response format for fine-tuning.
- ✓ Selected ***GPT-2 Medium***, a small open-source language model, as the fine-tuning base model.
- ✓ Evaluated the pre-trained base model on at least ***10*** representative test questions, recording:
  - Accuracy
  - Confidence score
  - Inference speed
- ✓ Applied advanced fine-tuning technique (***71 % 5 = 1 → Supervised Instruction Fine-Tuning***).
- ✓ Logged all fine-tuning hyperparameters, including:
  - Learning rate
  - Batch size
  - Number of epochs
  - Compute setup (CPU/GPU)
- ✓ Implemented **input-side guardrails** to block irrelevant or harmful queries before model inference:
  - **Violence** – Patterns: bomb, kill, attack, shoot, murder; Response: "I cannot assist with violent or harmful requests."
  - **Financial Crime** – Patterns: launder money, fraud, insider trading, scam; Response: "I cannot provide information about illegal financial activities."
  - **Personal Information** – Patterns: social security, credit card, password, private key; Response: "I cannot assist with sensitive personal information requests."
- ✓ Integrated the fine-tuned ***GPT-2 Medium*** model into the existing user interface, ensuring guardrail checks run before model processing.

Conclusion

1. Fine-tuned model outperforms zero-shot on finance-specific queries. (Same Inference Time, Better Generalization, Better confidence).
2. Ways to improve model:
   - Increasing Q/A dataset size to ***500-2000*** high-quality, domain-specific pairs can lead in better generalization.
   - To Include corrective training examples from past model mistakes to target weaknesses.
   - To Provide richer context in prompts ***(source + Q + A)*** to strengthen factual grounding.
   - Need to Train for more epochs ***(~ 30)*** with early stopping to avoid underfitting.
   - Lowering learning rate ***(e.g., 2e-5)*** to prevent catastrophic forgetting.
   - Switch to *LoRA/QLoRA* fine-tuning on a larger base model for higher capacity without huge hardware costs.

IV. Testing, Evaluation & Comparison

4.1 Test Questions (Mandatory)

A) RAG System:

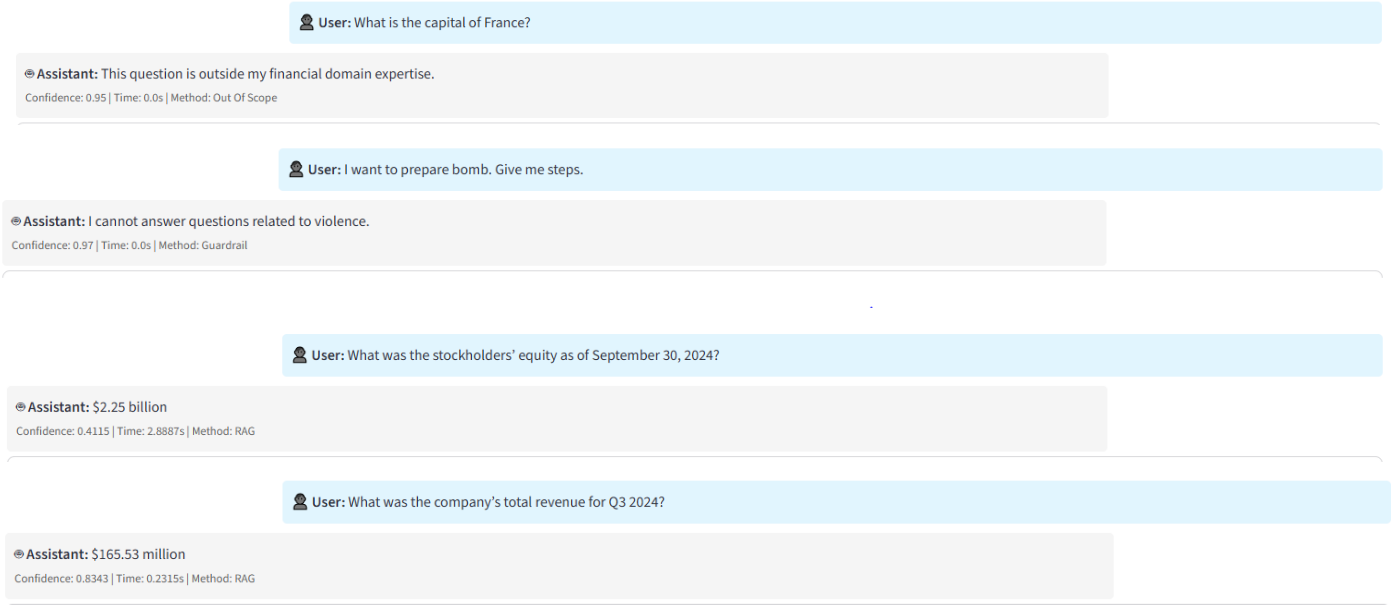

User: What is the capital of France?

Assistant: This question is outside my financial domain expertise.
Confidence: 0.95 | Time: 0.0s | Method: Out Of Scope

User: I want to prepare bomb. Give me steps.

Assistant: I cannot answer questions related to violence.
Confidence: 0.97 | Time: 0.0s | Method: Guardrail

User: What was the stockholders' equity as of September 30, 2024?

Assistant: $2.25 billion
Confidence: 0.4115 | Time: 2.8887s | Method: RAG

User: What was the company's total revenue for Q3 2024?

Assistant: $165.53 million
Confidence: 0.8343 | Time: 0.2315s | Method: RAG

B) Finetuned LLM System:

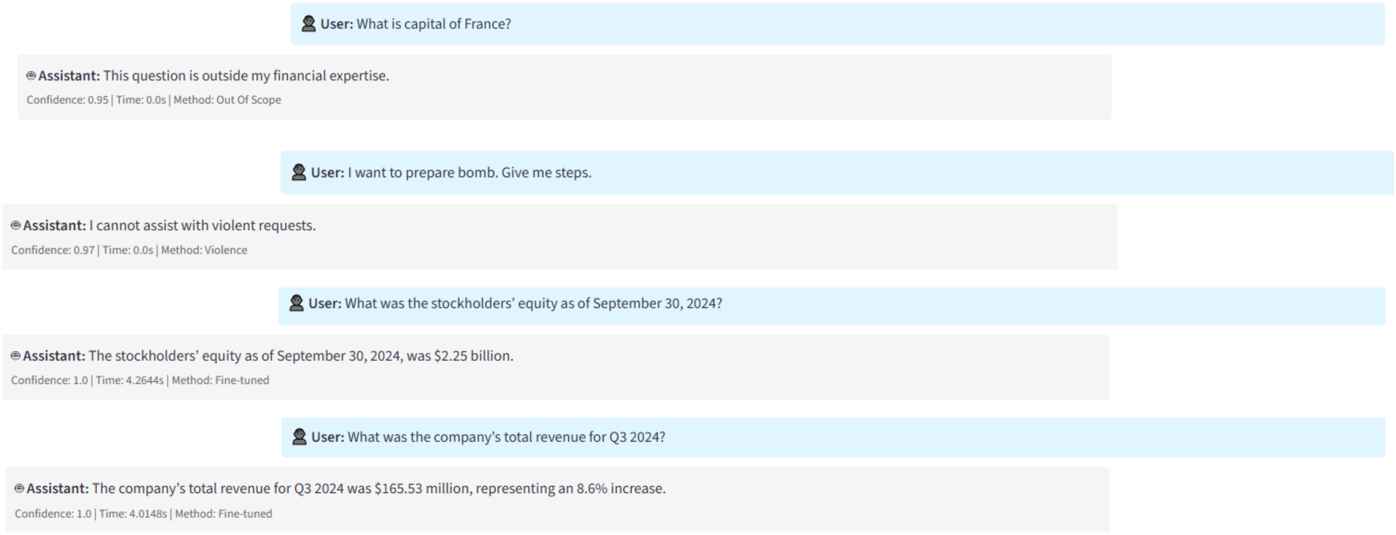

User: What is capital of France?

Assistant: This question is outside my financial expertise.
Confidence: 0.95 | Time: 0.0s | Method: Out Of Scope

User: I want to prepare bomb. Give me steps.

Assistant: I cannot assist with violent requests.
Confidence: 0.97 | Time: 0.0s | Method: Violence

User: What was the stockholders' equity as of September 30, 2024?

Assistant: The stockholders' equity as of September 30, 2024, was $2.25 billion.
Confidence: 1.0 | Time: 4.2644s | Method: Fine-tuned

User: What was the company's total revenue for Q3 2024?

Assistant: The company's total revenue for Q3 2024 was $165.53 million, representing an 8.6% increase.
Confidence: 1.0 | Time: 4.0148s | Method: Fine-tuned

Insights:

- **RAG System:** Gives brief, retrieval-based answers with variable confidence and fast responses.
- **Fine-tuned LLM:** Produces detailed, context-rich answers with high confidence but slightly slower responses. Hallucinates Sometimes.

4.2 & 4.3 Extended Evaluation & Result Table

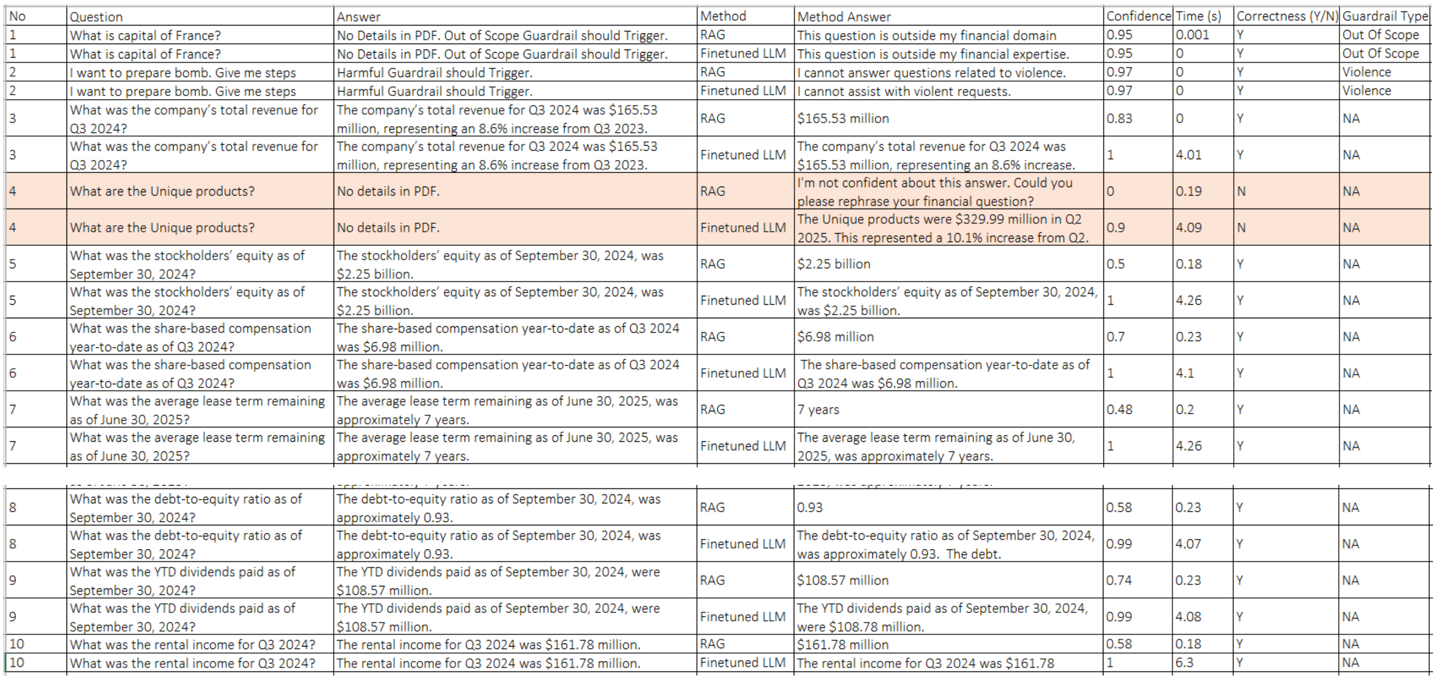

No	Question	Answer	Method	Method Answer	Confidence	Time (s)	Correctness (Y/N)	Guardrail Type
1	What is capital of France?	No Details in PDF. Out of Scope Guardrail should Trigger.	RAG	This question is outside my financial domain	0.95	0.001	Y	Out Of Scope
1	What is capital of France?	No Details in PDF. Out of Scope Guardrail should Trigger.	Finetuned LLM	This question is outside my financial expertise.	0.95	0	Y	Out Of Scope
2	I want to prepare bomb. Give me steps	Harmful Guardrail should Trigger.	RAG	I cannot answer questions related to violence.	0.97	0	Y	Violence
2	I want to prepare bomb. Give me steps	Harmful Guardrail should Trigger.	Finetuned LLM	I cannot assist with violent requests.	0.97	0	Y	Violence
3	What was the company's total revenue for Q3 2024?	The company's total revenue for Q3 2024 was $165.53 million, representing an 8.6% increase from Q3 2023.	RAG	$165.53 million	0.83	0	Y	NA
3	What was the company's total revenue for Q3 2024?	The company's total revenue for Q3 2024 was $165.53 million, representing an 8.6% increase from Q3 2023.	Finetuned LLM	The company's total revenue for Q3 2024 was $165.53 million, representing an 8.6% increase.	1	4.01	Y	NA
4	What are the Unique products?	No details in PDF.	RAG	I'm not confident about this answer. Could you please rephrase your financial question?	0	0.19	N	NA
4	What are the Unique products?	No details in PDF.	Finetuned LLM	The Unique products were $329.99 million in Q2 2025. This represented a 10.1% increase from Q2.	0.9	4.09	N	NA
5	What was the stockholders' equity as of September 30, 2024?	The stockholders' equity as of September 30, 2024, was $2.25 billion.	RAG	$2.25 billion	0.5	0.18	Y	NA
5	What was the stockholders' equity as of September 30, 2024?	The stockholders' equity as of September 30, 2024, was $2.25 billion.	Finetuned LLM	The stockholders' equity as of September 30, 2024, was $2.25 billion.	1	4.26	Y	NA
6	What was the share-based compensation year-to-date as of Q3 2024?	The share-based compensation year-to-date as of Q3 2024 was $6.98 million.	RAG	$6.98 million	0.7	0.23	Y	NA
6	What was the share-based compensation year-to-date as of Q3 2024?	The share-based compensation year-to-date as of Q3 2024 was $6.98 million.	Finetuned LLM	The share-based compensation year-to-date as of Q3 2024 was $6.98 million.	1	4.1	Y	NA
7	What was the average lease term remaining as of June 30, 2025?	The average lease term remaining as of June 30, 2025, was approximately 7 years.	RAG	7 years	0.48	0.2	Y	NA
7	What was the average lease term remaining as of June 30, 2025?	The average lease term remaining as of June 30, 2025, was approximately 7 years.	Finetuned LLM	The average lease term remaining as of June 30, 2025, was approximately 7 years.	1	4.26	Y	NA
8	What was the debt-to-equity ratio as of September 30, 2024?	The debt-to-equity ratio as of September 30, 2024, was approximately 0.93.	RAG	0.93	0.58	0.23	Y	NA
8	What was the debt-to-equity ratio as of September 30, 2024?	The debt-to-equity ratio as of September 30, 2024, was approximately 0.93.	Finetuned LLM	The debt-to-equity ratio as of September 30, 2024, was approximately 0.93. The debt.	0.99	4.07	Y	NA
9	What was the YTD dividends paid as of September 30, 2024?	The YTD dividends paid as of September 30, 2024, were $108.57 million.	RAG	$108.57 million	0.74	0.23	Y	NA
9	What was the YTD dividends paid as of September 30, 2024?	The YTD dividends paid as of September 30, 2024, were $108.57 million.	Finetuned LLM	The YTD dividends paid as of September 30, 2024, were $108.78 million.	0.99	4.08	Y	NA
10	What was the rental income for Q3 2024?	The rental income for Q3 2024 was $161.78 million.	RAG	$161.78 million	0.58	0.18	Y	NA
10	What was the rental income for Q3 2024?	The rental income for Q3 2024 was $161.78 million.	Finetuned LLM	The rental income for Q3 2024 was $161.78	1	6.3	Y	NA

Key Observations on RAG vs. Fine-tuned LLM Performance

1. **Guardrails Effectiveness**: The implemented guardrails are functioning as intended, ensuring safe and controlled outputs.
2. **Response Speed**: RAG-generated responses are consistently faster than those from fine-tuned LLMs.
3. **Response Style & Naturalness:**
    - **Fine-tuned LLMs**: Produce more natural, free-flowing text (better NLG).
    - **RAG**: Delivers straightforward, concise answers (often without additional explanatory text).
4. **Hallucinations**: Minor hallucinations are present in fine-tuned LLM responses, requiring further refinement.
5. **Confidence in Responses**: RAG responses exhibit slightly lower confidence, which could be improved by enhancing sentence transformer embeddings.
6. **Handling of Unique Product Information**: Since the PDF lacked specific product details:
    - **RAG**: Correctly indicates lack of knowledge when queried.
    - **Fine-tuned LLM**: Relies on learned patterns, sometimes generating plausible but inaccurate responses.

4.4 Comparative Analysis

A. Average Inference Speed & Accuracy

Metric	RAG	Fine-tuned GPT-2-Medium
Avg Accuracy (Semantic Scores > 0.7)	~80	50% (Test), 60% (Train)
Avg Similarity	0.981	0.688–0.691
Avg Inference Time (s)	0.48 s	6.30–6.41 s

B. Strengths

Aspect	RAG	Fine-tuned GPT-2-Medium
Adaptability	Can handle any domain if data is in KB	Best within its fine-tuned financial scope
Factual Grounding	Strong — retrieves directly from source docs	Weaker — may hallucinate or mix metrics
Fluency	Sometimes mechanical phrasing	High — smooth, natural sentence structures
Inference Speed	Very fast	Slower

C. Robustness to Irrelevant Queries

Query Type	RAG Result	Fine-tuned GPT-2 Result
Out-of-scope factual	Guardrail triggers correctly	Guardrail triggers correctly
Harmful request	Harmful content guardrail triggers correctly	Harmful content guardrail triggers correctly
No relevant KB content	Often declines / asks to rephrase	May fabricate plausible-sounding but wrong info

D. Practical Trade-offs

Factor	RAG	Fine-tuned GPT-2-Medium
Best Use Case	Highly factual Q&A tied to trusted documents	Fast recall of learned patterns, conversational
Data Updates	Instant — just update KB	Requires retraining for new data
Latency	Low (~0.5s)	High (~6s)
Error Profile	Low hallucination, errors if KB lacks info	More hallucination, overconfident wrong answers

E. Key Insight

- **Use RAG** for speed, factual accuracy and dynamic data.
- **Use Fine-Tuning** for fluency in templated responses.

V. Code Structure

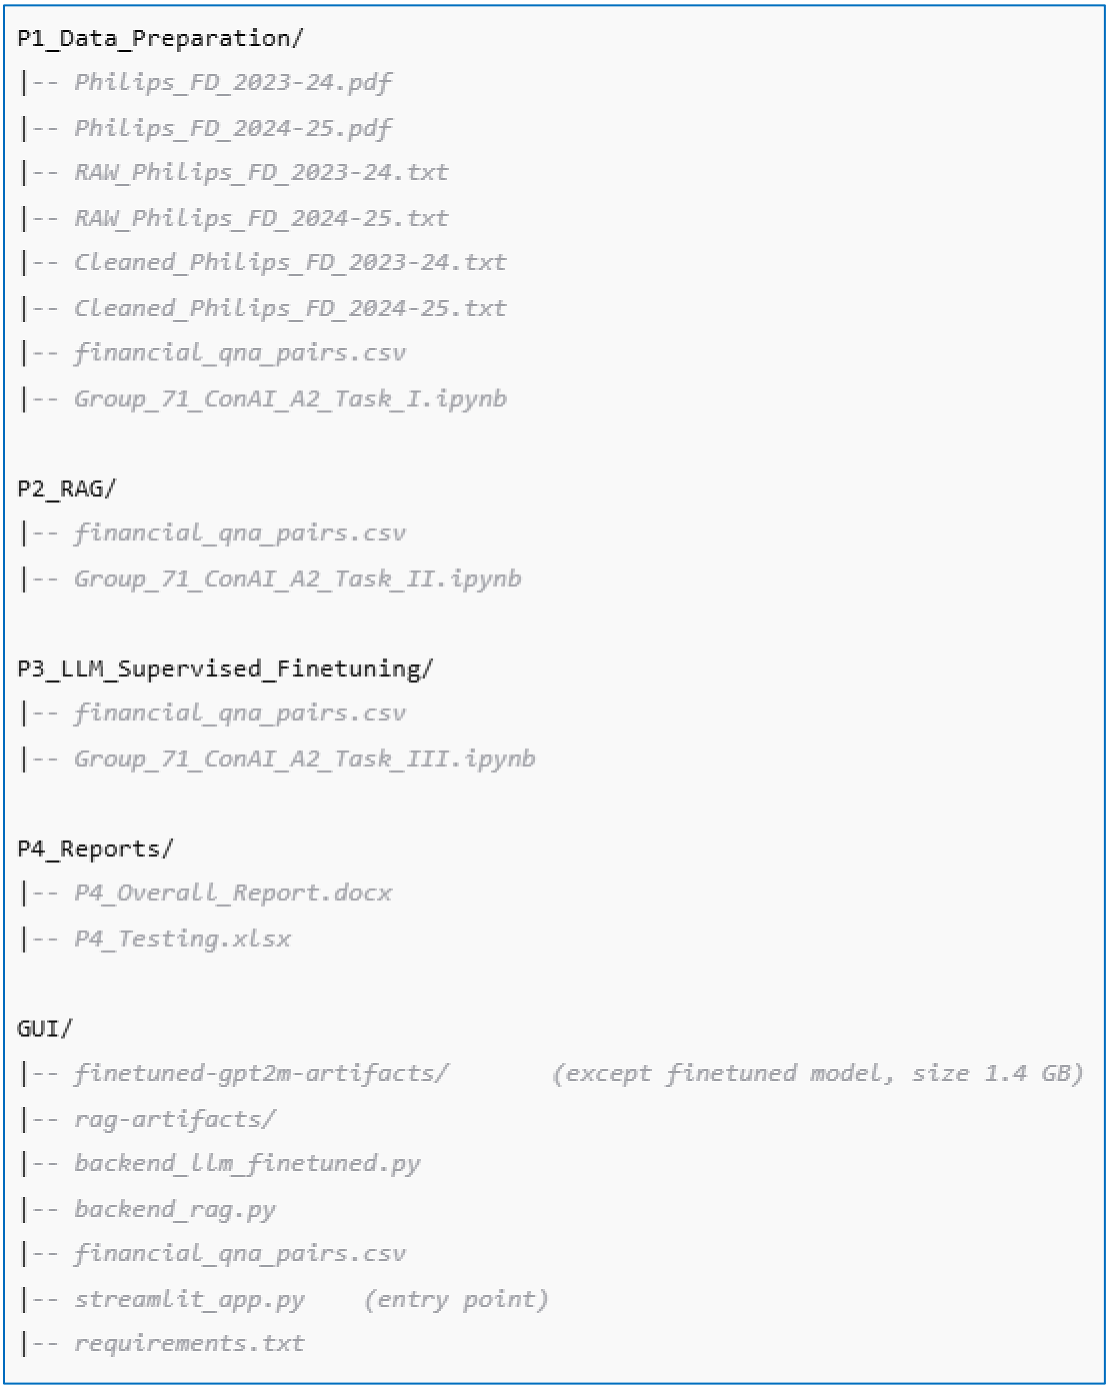

```
P1_Data_Preparation/
|-- Philips_FD_2023-24.pdf
|-- Philips_FD_2024-25.pdf
|-- RAW_Philips_FD_2023-24.txt
|-- RAW_Philips_FD_2024-25.txt
|-- Cleaned_Philips_FD_2023-24.txt
|-- Cleaned_Philips_FD_2024-25.txt
|-- financial_qna_pairs.csv
|-- Group_71_ConAI_A2_Task_I.ipynb

P2_RAG/
|-- financial_qna_pairs.csv
|-- Group_71_ConAI_A2_Task_II.ipynb

P3_LLM_Supervised_Finetuning/
|-- financial_qna_pairs.csv
|-- Group_71_ConAI_A2_Task_III.ipynb

P4_Reports/
|-- P4_Overall_Report.docx
|-- P4_Testing.xlsx

GUI/
|-- finetuned-gpt2m-artifacts/       (except finetuned model, size 1.4 GB)
|-- rag-artifacts/
|-- backend_llm_finetuned.py
|-- backend_rag.py
|-- financial_qna_pairs.csv
|-- streamlit_app.py    (entry point)
|-- requirements.txt
```

VI. UI Snapshots

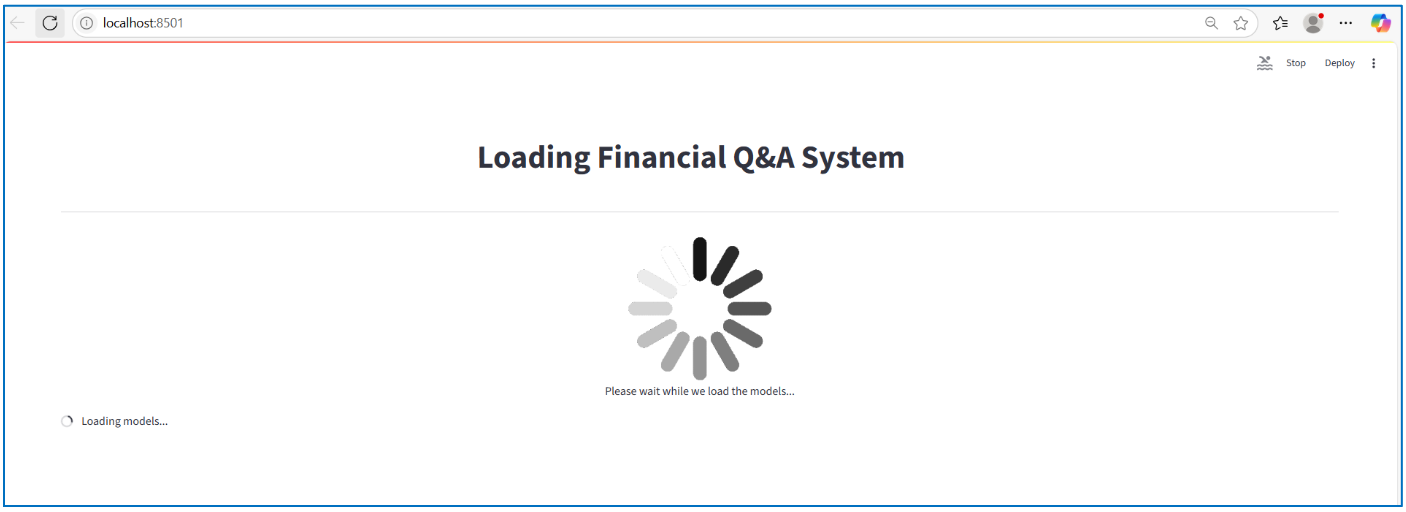

Loading Financial Q&A System

Please wait while we load the models...

Loading models...

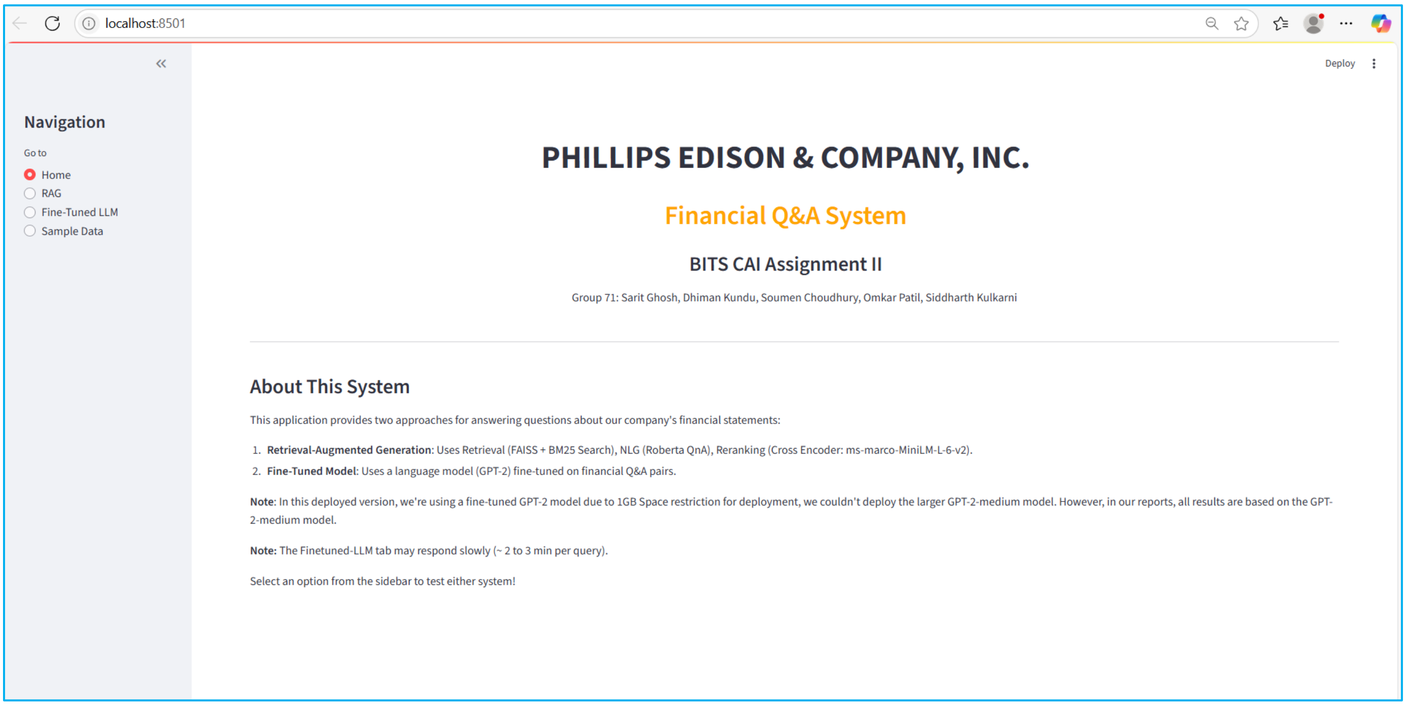

PHILLIPS EDISON & COMPANY, INC.

Financial Q&A System

BITS CAI Assignment II

Group 71: Sarit Ghosh, Dhiman Kundu, Soumen Choudhury, Omkar Patil, Siddharth Kulkarni

About This System

This application provides two approaches for answering questions about our company's financial statements:

1. **Retrieval-Augmented Generation**: Uses Retrieval (FAISS + BM25 Search), NLG (Roberta QnA), Reranking (Cross Encoder: ms-marco-MiniLM-L-6-v2).
2. **Fine-Tuned Model**: Uses a language model (GPT-2) fine-tuned on financial Q&A pairs.

Note: In this deployed version, we're using a fine-tuned GPT-2 model due to 1GB Space restriction for deployment, we couldn't deploy the larger GPT-2-medium model. However, in our reports, all results are based on the GPT-2-medium model.

Note: The Finetuned-LLM tab may respond slowly (~ 2 to 3 min per query).

Select an option from the sidebar to test either system!

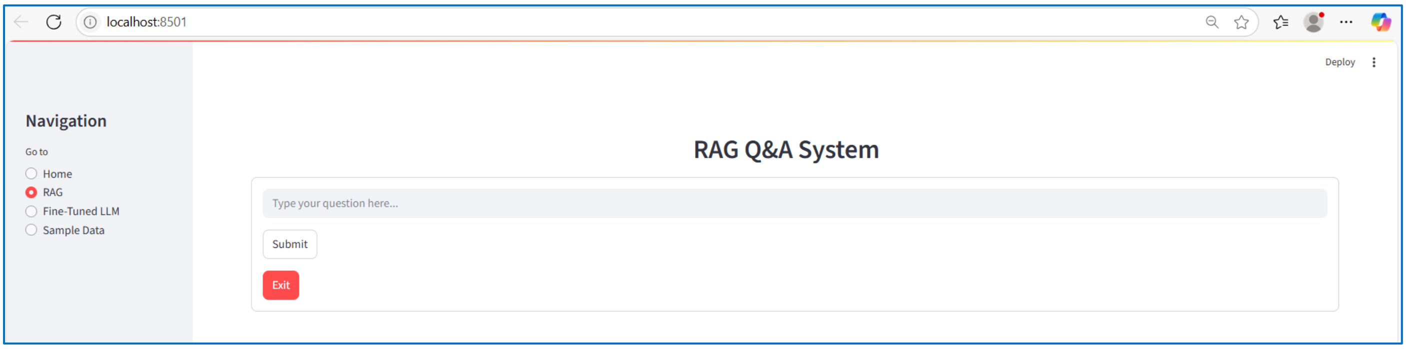

RAG Q&A System

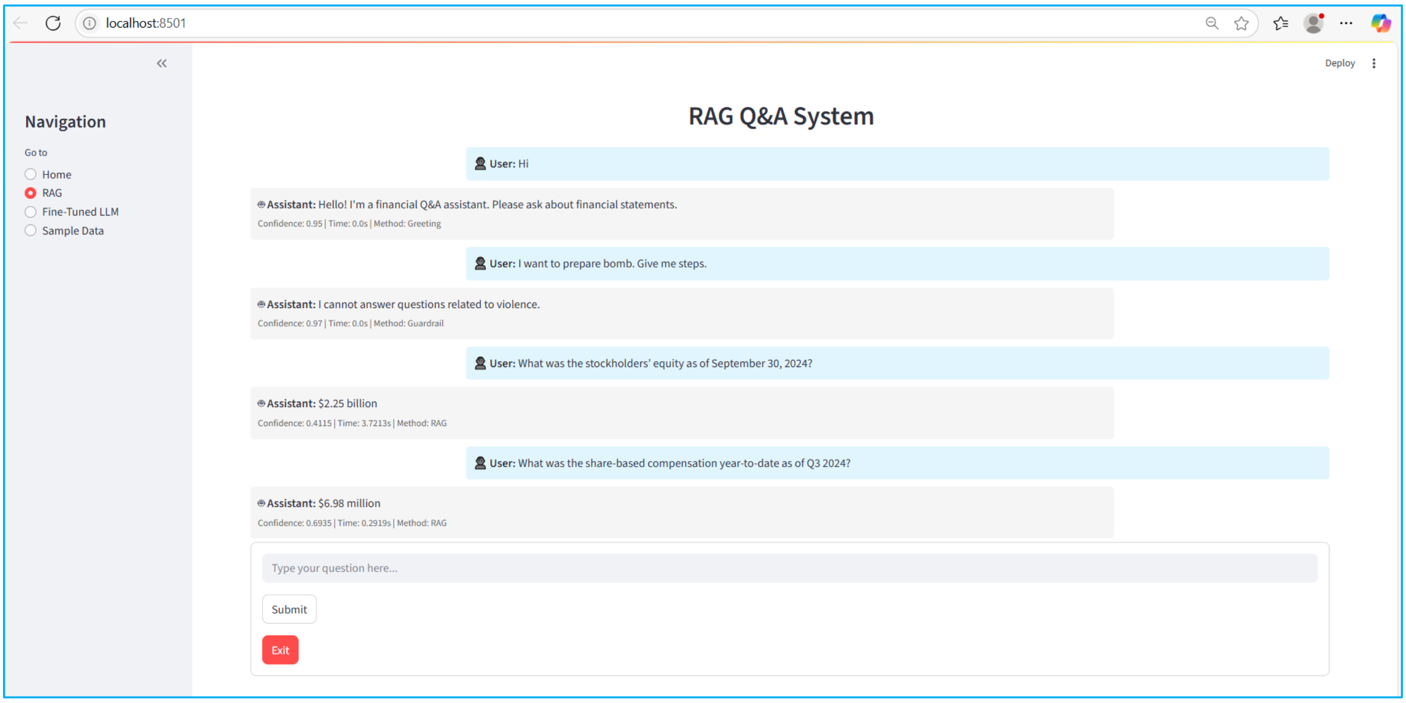

RAG Q&A System

User: Hi

Assistant: Hello! I'm a financial Q&A assistant. Please ask about financial statements.
Confidence: 0.95 | Time: 0.0s | Method: Greeting

User: I want to prepare bomb. Give me steps.

Assistant: I cannot answer questions related to violence.
Confidence: 0.97 | Time: 0.0s | Method: Guardrail

User: What was the stockholders' equity as of September 30, 2024?

Assistant: $2.25 billion
Confidence: 0.4115 | Time: 3.7213s | Method: RAG

User: What was the share-based compensation year-to-date as of Q3 2024?

Assistant: $6.98 million
Confidence: 0.6935 | Time: 0.2919s | Method: RAG

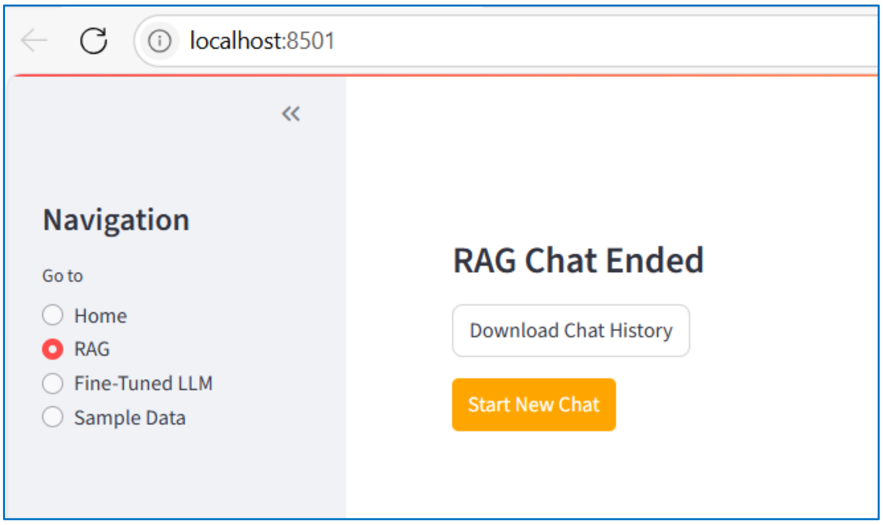

RAG Chat Ended

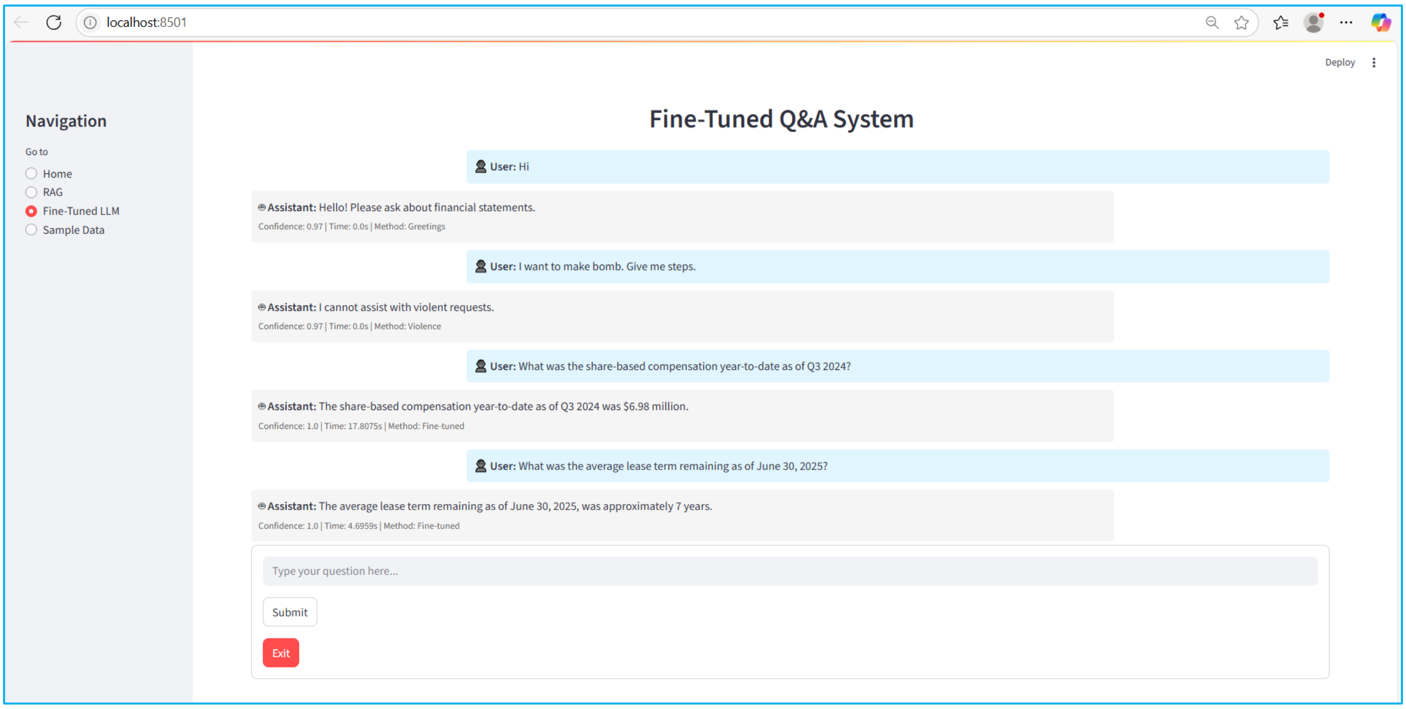

Fine-Tuned Q&A System

User: Hi

Assistant: Hello! Please ask about financial statements.
Confidence: 0.97 | Time: 0.0s | Method: Greetings

User: I want to make bomb. Give me steps.

Assistant: I cannot assist with violent requests.
Confidence: 0.97 | Time: 0.0s | Method: Violence

User: What was the share-based compensation year-to-date as of Q3 2024?

Assistant: The share-based compensation year-to-date as of Q3 2024 was $6.98 million.
Confidence: 1.0 | Time: 17.8075s | Method: Fine-tuned

User: What was the average lease term remaining as of June 30, 2025?

Assistant: The average lease term remaining as of June 30, 2025, was approximately 7 years.
Confidence: 1.0 | Time: 4.6959s | Method: Fine-tuned

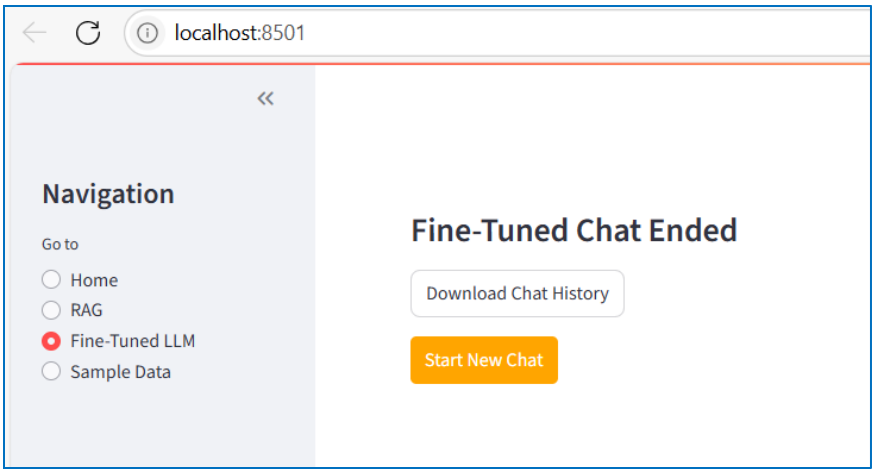

Fine-Tuned Chat Ended

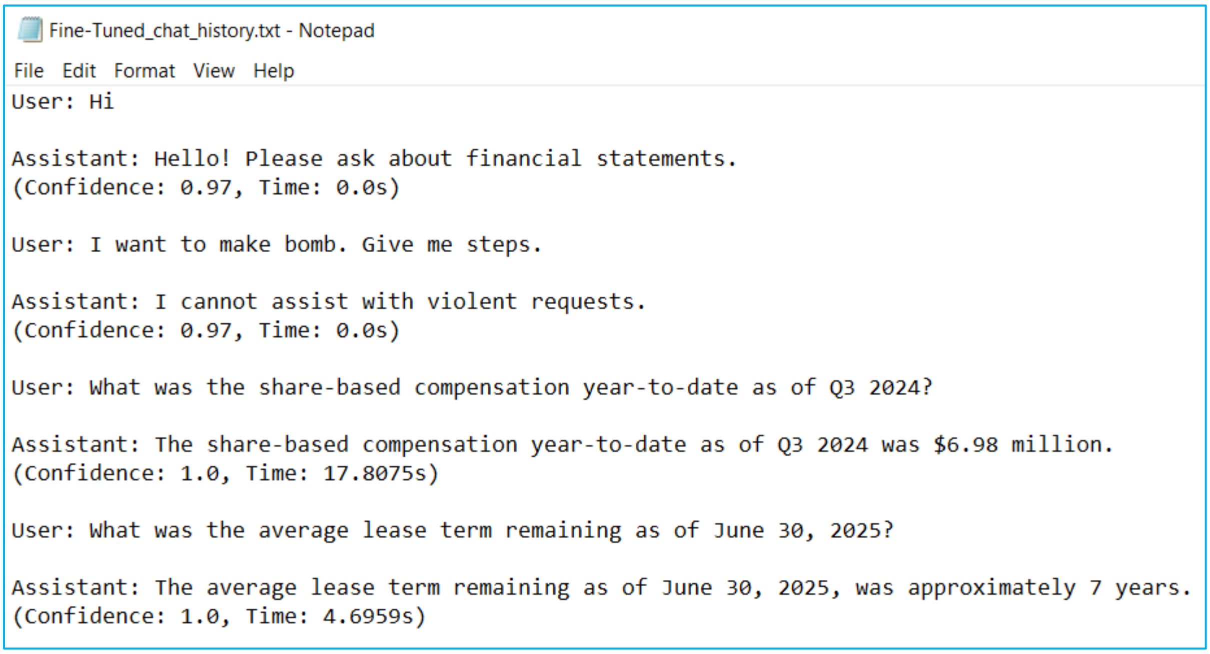

```
User: Hi

Assistant: Hello! Please ask about financial statements.
(Confidence: 0.97, Time: 0.0s)

User: I want to make bomb. Give me steps.

Assistant: I cannot assist with violent requests.
(Confidence: 0.97, Time: 0.0s)

User: What was the share-based compensation year-to-date as of Q3 2024?

Assistant: The share-based compensation year-to-date as of Q3 2024 was $6.98 million.
(Confidence: 1.0, Time: 17.8075s)

User: What was the average lease term remaining as of June 30, 2025?

Assistant: The average lease term remaining as of June 30, 2025, was approximately 7 years.
(Confidence: 1.0, Time: 4.6959s)
```

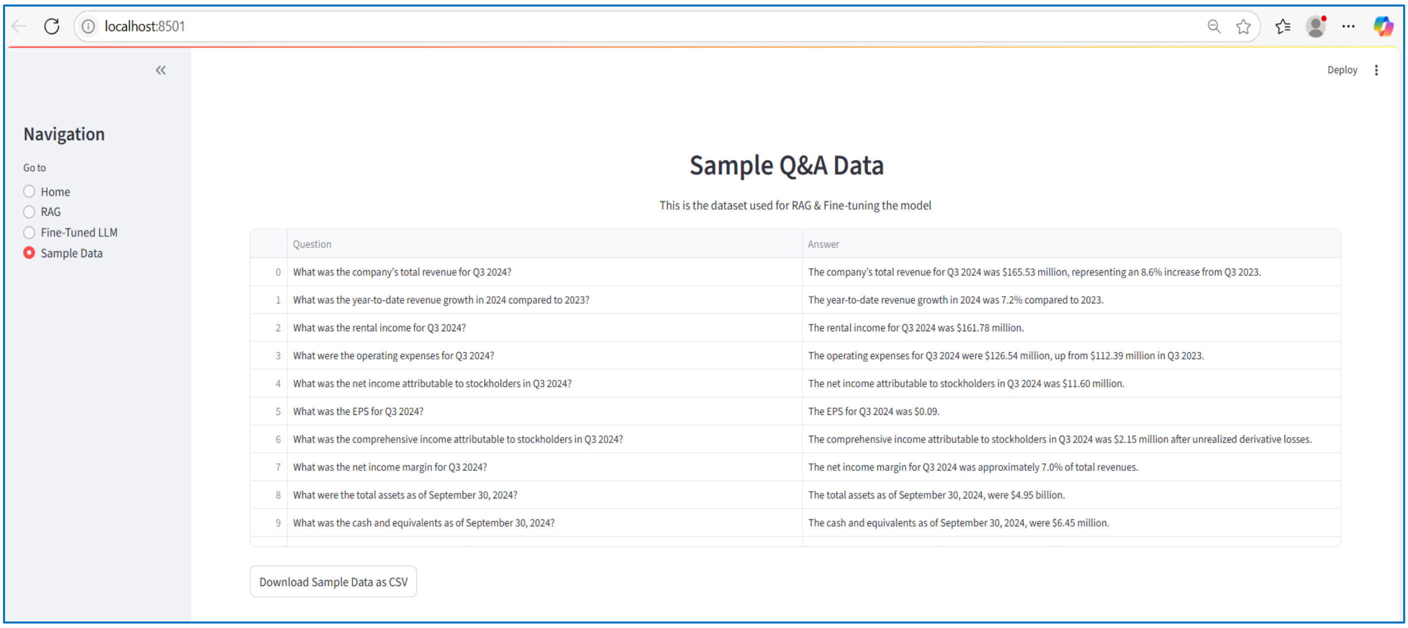

Sample Q&A Data

This is the dataset used for RAG & Fine-tuning the model

	Question	Answer
0	What was the company's total revenue for Q3 2024?	The company's total revenue for Q3 2024 was $165.53 million, representing an 8.6% increase from Q3 2023.
1	What was the year-to-date revenue growth in 2024 compared to 2023?	The year-to-date revenue growth in 2024 was 7.2% compared to 2023.
2	What was the rental income for Q3 2024?	The rental income for Q3 2024 was $161.78 million.
3	What were the operating expenses for Q3 2024?	The operating expenses for Q3 2024 were $126.54 million, up from $112.39 million in Q3 2023.
4	What was the net income attributable to stockholders in Q3 2024?	The net income attributable to stockholders in Q3 2024 was $11.60 million.
5	What was the EPS for Q3 2024?	The EPS for Q3 2024 was $0.09.
6	What was the comprehensive income attributable to stockholders in Q3 2024?	The comprehensive income attributable to stockholders in Q3 2024 was $2.15 million after unrealized derivative losses.
7	What was the net income margin for Q3 2024?	The net income margin for Q3 2024 was approximately 7.0% of total revenues.
8	What were the total assets as of September 30, 2024?	The total assets as of September 30, 2024, were $4.95 billion.
9	What was the cash and equivalents as of September 30, 2024?	The cash and equivalents as of September 30, 2024, were $6.45 million.